

Google PageRank: The Random Surfer, Markov Chains, Convergence, and Scalable Computation

Varnit Sahu

Abstract

PageRank converts the hyperlink structure of the Web into a numerical measure of importance. Its central idea is simple: an important page is one that is visited frequently by a hypothetical random surfer, and a link from an important page carries more weight than a link from an unimportant one. Turning this circular definition into a reliable algorithm requires several ideas from numerical linear algebra: stochastic matrices, stationary distributions, the Perron–Frobenius theorem, and the power method. This article develops PageRank from the random-surfer model, treats dangling pages and rank traps carefully, proves existence, uniqueness, and convergence, analyzes time and storage costs, and then extends the basic model to personalization, topic-sensitive ranking, weighted links, trust and spam resistance, and dynamic graphs.

Problem

Suppose the Web contains n pages joined by directed hyperlinks. A search engine must answer two different questions:

- (1) Which pages are relevant to the words in a query?
- (2) Among the relevant pages, which pages are globally important?

PageRank addresses the second question. Unlike a simple count of incoming links, it assigns more influence to a link originating at an important page. This creates a recursive definition: a page is important when important pages link to it. Our goal is to make that definition mathematically precise and computationally useful.

The Web as a Directed Graph

Let $G = (V, E)$ be a directed graph with pages $V = \{1, \dots, n\}$. An edge $(j, i) \in E$ means page j contains a link to page i . Write

$$\deg^+(j) = \#\{i : (j, i) \in E\}$$

for the outdegree of page j . If page j divides its recommendation equally among its outlinks, then page i receives a fraction $1/\deg^+(j)$ of j 's current importance.

A naive score such as the indegree $\deg^-(i)$ treats all incoming links equally. PageRank instead seeks scores p_i satisfying the recursive relation

$$p_i \approx \sum_{j \rightarrow i} \frac{p_j}{\deg^+(j)}.$$

This is already an eigenvector equation, but it is not yet safe: pages with no outlinks and closed cycles can prevent a unique, meaningful solution.

The Random-Surfer Model

Imagine a surfer located at page j . At each step the surfer chooses one of j 's outgoing links uniformly and follows it. After many steps, the fraction of time spent on page i estimates the importance of that page.

A literal simulation can start many walks and count visits. If K steps are taken from each of n starting pages, the work is $\Theta(nK)$. More importantly, a finite random sample introduces variance: obtaining an additional decimal digit of Monte Carlo accuracy usually requires about one hundred times as many samples. PageRank replaces this sampling process by deterministic probability propagation.

From the Surfer to a Transition Matrix

Define the hyperlink matrix $P \in \mathbb{R}^{n \times n}$ by

$$P_{ij} = \begin{cases} 1/\deg^+(j), & \text{if } j \rightarrow i, \\ 0, & \text{otherwise.} \end{cases}$$

We use the column-stochastic convention: column j describes where the surfer moves from page j . If $p^{(k)}$ is the probability distribution after k steps, then

$$p^{(k+1)} = Pp^{(k)}.$$

Definition 1. A vector $p \in \mathbb{R}^n$ is a probability vector if $p_i \geq 0$ and $\mathbf{e}^T p = 1$, where $\mathbf{e} = (1, \dots, 1)^T$. A matrix S is column stochastic if $S_{ij} \geq 0$ and $\mathbf{e}^T S = \mathbf{e}^T$.

Lemma 1 (Conservation of probability). If S is column stochastic and p is a probability vector, then Sp is a probability vector.

Proof. Nonnegativity follows from $S \geq 0$ and $p \geq 0$. Moreover,

$$\mathbf{e}^T (Sp) = (\mathbf{e}^T S)p = \mathbf{e}^T p = 1.$$

□

Two Structural Failures

Dangling pages

A dangling page has no outlinks. Its column in P is zero, so probability disappears when the surfer reaches it. Let

$$d_j = \begin{cases} 1, & \deg^+(j) = 0, \\ 0, & \text{otherwise,} \end{cases}$$

and let v be any probability vector. Replacing every dangling column by v gives

$$S = P + vd^T.$$

For uniform redistribution, $v = \mathbf{e}/n$, which recovers the construction in the course notes. The matrix S is column stochastic.

Rank sinks, cycles, and reducibility

A group of pages may link only within the group. Once the surfer enters, it cannot leave. Periodic structures create another problem: on a two-page directed cycle, the probability vector can oscillate forever. Thus stochasticity alone does not guarantee a unique limit independent of the starting vector.

Teleportation and the Google Matrix

At each step, let the surfer

- follow a hyperlink with probability α , and
- teleport to a page sampled from v with probability $1 - \alpha$,

where $0 < \alpha < 1$. The resulting Google matrix is

$$M = \alpha S + (1 - \alpha)v\mathbf{e}^T.$$

Its j th column is $\alpha S_{:j} + (1 - \alpha)v$. Consequently every column sums to one. If every $v_i > 0$, then every entry of M is positive.

The parameter α controls a fundamental tradeoff. Large α follows the link graph faithfully but converges more slowly and is more sensitive to link manipulation. Small α converges quickly but pushes the result toward v . The familiar illustrative choice is $\alpha = 0.85$, although the mathematics does not require this value.

The PageRank Equation

Definition 2 (PageRank). *The PageRank vector is a probability vector p satisfying*

$$p = Mp.$$

Equivalently,

$$p = \alpha Sp + (1 - \alpha)v. \tag{1}$$

The score p_i is the long-run probability that the random surfer is at page i .

Rearranging (1) gives the linear system

$$(I - \alpha S)p = (1 - \alpha)v. \tag{2}$$

Because every eigenvalue of a stochastic matrix has magnitude at most one and $\alpha < 1$, $I - \alpha S$ is nonsingular. Hence (2) has a unique solution. At Web scale, however, a sparse iterative method is preferable to a general direct factorization.

A Small Worked Example

Consider four pages with links

$$1 \rightarrow \{2, 3\}, \quad 2 \rightarrow 3, \quad 3 \rightarrow 1, \quad 4 \text{ dangling}.$$

With uniform dangling redistribution and uniform teleportation,

$$P = \begin{bmatrix} 0 & 0 & 1 & 0 \\ 1/2 & 0 & 0 & 0 \\ 1/2 & 1 & 0 & 0 \\ 0 & 0 & 0 & 0 \end{bmatrix}, \quad d = \begin{bmatrix} 0 \\ 0 \\ 0 \\ 1 \end{bmatrix}, \quad v = \frac{1}{4}\mathbf{e},$$

and $S = P + vd^T$. For $\alpha = 0.85$, the update can be evaluated without explicitly constructing the dense matrix M :

$$p^{(k+1)} = 0.85 \left(Pp^{(k)} + v d^T p^{(k)} \right) + 0.15v.$$

Starting from $p^{(0)} = v$, the iteration converges approximately to

$$p \approx (0.3693, 0.2046, 0.3785, 0.0476)^T.$$

Page 3 ranks first, page 1 is close behind, and the dangling page ranks last. Notice that page 1 has only one incoming link, but that link comes from the highly ranked page 3; PageRank captures this recursive reinforcement.

Existence, Uniqueness, and Convergence

Theorem 1 (Positive-matrix PageRank). *Assume $0 < \alpha < 1$ and $v_i > 0$ for every i . Then M has a unique stationary probability vector p . For every initial probability vector $p^{(0)}$,*

$$M^k p^{(0)} \longrightarrow p.$$

Proof. The matrix M is positive and column stochastic. The Perron–Frobenius theorem therefore gives a simple dominant eigenvalue equal to its spectral radius, which is 1, and an associated strictly positive eigenvector. Normalizing that eigenvector to have unit sum gives the unique stationary probability vector. Every other eigenvalue has magnitude strictly less than 1, so all nonstationary components of $M^k p^{(0)}$ decay to zero. $\square$

There is also a direct contraction argument that gives a useful error bound. For any two probability vectors x and y , $\mathbf{e}^T(x - y) = 0$ and

$$M(x - y) = \alpha S(x - y).$$

Since a column-stochastic matrix is nonexpansive in the 1-norm,

$$\|Mx - My\|_1 \leq \alpha \|x - y\|_1. \quad (3)$$

Thus the PageRank map is a contraction with factor at most α .

Proposition 1 (A posteriori stopping bound). *If p is the PageRank vector, then*

$$\left\| p - p^{(k)} \right\|_1 \leq \frac{\alpha}{1 - \alpha} \left\| p^{(k)} - p^{(k-1)} \right\|_1.$$

Proof. Apply (3) repeatedly to the telescoping tail $p - p^{(k)} = \sum_{r=k}^{\infty} (p^{(r+1)} - p^{(r)})$ and sum the resulting geometric series. $\square$

Therefore, to guarantee $\left\| p - p^{(k)} \right\|_1 \leq \varepsilon$, it suffices to stop when

$$\left\| p^{(k)} - p^{(k-1)} \right\|_1 \leq \frac{1 - \alpha}{\alpha} \varepsilon.$$

The residual $r^{(k)} = \alpha S p^{(k)} + (1 - \alpha)v - p^{(k)}$ provides a similar computable test.

The Power Method

The basic algorithm is backward-free and memory efficient:

1. Choose a probability vector $p^{(0)}$ (often v).
2. Compute the dangling mass $\delta_k = d^T p^{(k)}$.
3. Update

$$p^{(k+1)} = \alpha(Pp^{(k)} + v\delta_k) + (1 - \alpha)v.$$

4. Stop when the change or residual is below tolerance.

In exact arithmetic each iterate is a probability vector. In floating-point arithmetic, small negative values or sum drift can be corrected by clipping tiny roundoff errors and renormalizing, though a stable implementation should require little correction.

Complexity Analysis

Let $m = |E|$ be the number of hyperlinks and let k be the number of iterations.

Why the dense matrix must not be formed

The teleportation matrix ve^T is dense. Explicitly storing M costs $\Theta(n^2)$ memory and a dense matrix-vector multiplication costs $\Theta(n^2)$ time. This is infeasible even for moderately large graphs.

The formula above keeps only the sparse hyperlink matrix P , the vectors p, v, d , and a scalar dangling mass. In compressed sparse-column or adjacency-list form:

- sparse multiplication $Pp^{(k)}$ costs $\Theta(m)$;
- dangling-mass computation and vector updates cost $\Theta(n)$;
- one iteration costs $\Theta(m + n)$;
- k iterations cost $\Theta(k(m + n))$;
- storage costs $\Theta(m + n)$.

When the average outdegree is bounded, $m = \Theta(n)$ and the per-iteration cost is linear in the number of pages. The more precise bound is $\Theta(m + n)$, not simply $\Theta(n)$ for an arbitrary graph.

Iteration count

From the contraction bound, reducing an initial error by a factor ε requires approximately

$$k \geq \frac{\log \varepsilon}{\log \alpha}$$

iterations in the worst-case geometric model. With $\alpha = 0.85$, $0.85^{114} \approx 9 \times 10^{-9}$, matching the scale discussed in the course notes. Actual convergence can be faster when the subdominant eigenvalues are substantially smaller than α .

Monte Carlo versus power iteration

Random-walk simulation can be distributed easily and can answer localized questions without storing a full rank vector, but its statistical error decays slowly. Power iteration is deterministic, exploits sparse linear algebra, and computes all scores simultaneously. Monte Carlo becomes attractive when only a few personalized scores are needed or when graph access is naturally stream-based.

Implementation at Scale

Several engineering choices preserve the mathematics while improving throughput:

- **Sparse graph layout.** Store destinations grouped by source so each source score is divided once and scattered along its outgoing edges.
- **Partitioning.** Distribute pages and edges across workers. Each iteration requires communication of contributions that cross partitions, followed by a global sum for dangling mass and convergence statistics.
- **Combining contributions.** Locally aggregate messages with the same destination to reduce network traffic.
- **Warm starts.** For a slowly changing graph, initialize with the previous PageRank vector rather than v .
- **Checkpointing.** Persist periodic iterates so a long distributed run can recover after worker failure.
- **Top- K stability.** If only ordering matters, monitor whether the leading set and score gaps have stabilized, while retaining a norm-based residual as the correctness check.

Extensions and Additional Features

Personalized and topic-sensitive PageRank

Uniform teleportation uses $v = \mathbf{e}/n$. A nonuniform v concentrates teleportation on trusted, topical, geographic, or user-preferred pages. Several topic vectors $v^{(1)}, \dots, v^{(q)}$ can be precomputed; at query time their PageRank vectors can be combined using topic weights. Because (2) is linear in v , if $v = \sum_{\ell} w_{\ell} v^{(\ell)}$, then

$$p(v) = \sum_{\ell} w_{\ell} p(v^{(\ell)}).$$

Weighted hyperlinks

Equal division among outlinks is a modeling choice, not a requirement. Links can be weighted by placement, freshness, semantic relevance, or estimated probability of a click. If $w_{ij} \geq 0$ measures the strength of link $j \rightarrow i$, define

$$P_{ij} = \frac{w_{ij}}{\sum_r w_{rj}}$$

for nondangling columns. All stochastic-matrix arguments continue to apply.

Trust, spam resistance, and fairness

Link farms attempt to manufacture dense mutual-link structures. Teleportation limits but does not eliminate this manipulation. A trust-biased vector can seed mass on vetted pages, while suspicious edges can be downweighted. These interventions introduce policy choices: the seed set, weighting rules, and crawl coverage can amplify established sites or disadvantage new and underlinked communities. PageRank should therefore be treated as one feature in a ranking system, not as an objective definition of quality.

Temporal PageRank

The Web changes continuously. Freshness can be incorporated by decaying old edges, emphasizing recent interactions, or updating only regions affected by graph changes. Warm-started power iteration is simple and robust; more advanced incremental methods estimate the correction to p after a localized update.

Multiple edge types

The same framework applies beyond Web links: citations, social follows, purchases, or transitions between application screens. For relation-specific stochastic matrices $S_1, \dots, S_r$ and weights $\beta_\ell \geq 0$ summing to one, one may use

$$S = \sum_{\ell=1}^r \beta_\ell S_\ell.$$

The resulting score blends several notions of endorsement while retaining the PageRank machinery.

Reverse and bipartite variants

Reversing edges highlights pages that point toward important destinations rather than pages endorsed by them. On bipartite graphs, alternating updates can rank two types of objects—for example, hubs and authorities, users and products, or authors and papers. These variants change the interpretation, so the edge direction and normalization convention must always be stated explicitly.

Sensitivity and the Role of α

Let $p(v)$ denote PageRank for teleportation vector v . From (2),

$$p(v) = (1 - \alpha)(I - \alpha S)^{-1}v.$$

The Neumann expansion

$$(I - \alpha S)^{-1} = \sum_{k=0}^{\infty} \alpha^k S^k$$

shows that PageRank is a weighted mixture of walks of every length. A path of length k receives weight proportional to α^k : larger α gives long link paths more influence.

For two teleportation vectors v and w , stochastic nonexpansiveness yields

$$\|p(v) - p(w)\|_1 \leq (1 - \alpha) \sum_{k=0}^{\infty} \alpha^k \|v - w\|_1 = \|v - w\|_1.$$

Thus personalization does not magnify 1-norm differences in its input. By contrast, as α approaches one, the linear system becomes more poorly conditioned in directions associated with slowly mixing graph structure, explaining slower convergence and greater sensitivity to graph perturbations.

Common Pitfalls

- **Mixing row and column conventions.** With column-stochastic matrices use $p^{(k+1)} = Mp^{(k)}$; with row-stochastic matrices use $(p^{(k+1)})^T = (p^{(k)})^T M$.
- **Ignoring dangling pages.** Zero columns destroy conservation of probability.
- **Forming the dense Google matrix.** Teleportation must be applied as a vector operation.
- **Using only a fixed iteration count.** A residual or difference test makes the achieved accuracy observable.
- **Equating PageRank with query relevance.** PageRank measures graph-based importance; a search system still needs textual, semantic, freshness, quality, and safety signals.
- **Treating scores as timeless or neutral.** Crawl boundaries, duplicate handling, edge weights, teleportation, and update frequency all affect the result.

Conclusion

PageRank turns a recursive social idea—important pages are endorsed by important pages—into a stationary-distribution problem. Dangling redistribution makes the hyperlink transition stochastic, while teleportation removes traps and produces a unique, globally attractive solution. The power method then computes that solution using only sparse matrix-vector products, with $\Theta(m+n)$ work per iteration and a convergence rate controlled by the damping factor and spectral gap.

The same foundation supports much more than a single global ranking. Personalization changes where the surfer teleports; weighted and multi-relational edges change what constitutes an endorsement; warm starts and distributed sparse operations address changing, massive graphs; and trust-aware variants respond to adversarial links. The enduring value of PageRank is therefore not merely one formula, but a reusable pattern: model movement on a graph, repair pathological structure, identify a stationary distribution, and compute it with scalable numerical linear algebra.

References

- [1] L. Page, S. Brin, R. Motwani, and T. Winograd, “The PageRank Citation Ranking: Bringing Order to the Web,” Stanford InfoLab, 1999.
- [2] A. N. Langville and C. D. Meyer, *Google’s PageRank and Beyond: The Science of Search Engine Rankings*, Princeton University Press, 2006.